

Do flood susceptibility models know more than the flood record?

Evidence from radar-observed floods in Bangladesh

Md Rakib Hasan^{1,2,*} 

¹Department of Soil, Water and Environment, University of Dhaka, Dhaka-1000, Bangladesh

²Fermium Systems, Dhaka-1207, Bangladesh

*Correspondence: rakibhridoy@fermium.systems

Running title: Flood records and susceptibility models

Abstract

Flood susceptibility maps are judged by how well they rank the ground that floods, yet they are rarely compared with the simplest alternative, the record of where floods have already been. This study makes that comparison for infrastructure mapped in OpenStreetMap in three regions of Bangladesh, with flood extents mapped from Sentinel-1 radar for thirteen events as labels and reference and every model scored on 10 km blocks held out from training. Trained on floods up to 2022, a gradient-boosted model ranked the ground the 2024 river floods reached in Rangpur and Rajshahi at an AUC of 0.801, below the record of earlier flooding at 0.896, and the record kept its lead with masks re-mapped from a single satellite, with two other training targets and with 40 km blocks. Given the record as a feature, the model beat both, at 0.924, by ranking the ground no earlier flood had reached. In the Sylhet haor the record's lead, 0.851 against 0.797, closed when the model learned the same target as the test, and on the cyclone-struck coast, where most flooded assets stood on ground no earlier cyclone had reached, the model's lead did not survive wider blocks. The record was therefore the stronger guide where floods returned to the same ground. Radar-observed extents also made better training labels than terrain thresholds, by 0.101, 0.087 and 0.080 in AUC on held-out areas, a graph neural network added nothing over tabular models, and elevation carried almost all of the terrain signal. The radar masks agreed only moderately with an independent MODIS map, so every score measures agreement with flooding as radar sees it. Susceptibility studies with a radar flood history should report that record as a baseline.

Keywords: flood susceptibility; Sentinel-1; spatial cross-validation; temporal validation; gradient boosting; graph neural network; Bangladesh

1 Introduction

Riverine and flash floods recur across northern Bangladesh in most monsoon seasons and cyclone surges strike its south-west coast, yet the susceptibility maps meant to show where flooding would do most harm are rarely tested against floods that actually occurred. The studies closest to this one report best-model AUCs between 0.87 and 0.98 (Islam et al., 2021; Talukdar et al., 2020; Islam and Chowdhury, 2024; Adnan et al., 2023), and two choices decide what such a figure means. The first is where the flood locations come from, usually a few hundred to a few thousand points drawn from field surveys, local knowledge or satellite flood maps. The second is how the test set is separated from the training set, since a random split of spatially clustered points tests a model mostly beside its own training points, a design known to inflate accuracy (Roberts et al., 2017; Ploton et al., 2020).

Both choices can now be made differently, because Sentinel-1 radar sees through monsoon cloud and has been processed into flood maps with overall accuracies above 94 % (Twele et al., 2016; Uddin et al., 2019). Every pixel observed to flood can therefore serve as a training label (Islam and Chowdhury, 2024), and whole areas can be held out so that the test asks how well a model carries to ground it has not seen. A radar

record of several floods also offers a baseline that susceptibility studies rarely report, which is the map of where water has already been. Graph neural networks, which pass information between neighbouring locations, have meanwhile entered hazard mapping (Wang et al., 2023a; Ma et al., 2025) on the appeal that a road or embankment floods together with its neighbours, although whether that structure adds anything over an ordinary tabular model has seldom been tested on identical data.

This study asks three questions of flood susceptibility models for mapped infrastructure in three regions with contrasting regimes, riverine flooding in Rangpur and Rajshahi, flash flooding in the Sylhet haor basin and cyclone surge on the south-west coast. The first, and the most consequential for how published maps should be read, is whether a model anticipates a later flood better than the record of earlier ones. The second is whether radar-observed flood extents make better training labels than terrain thresholds for ground a model has not seen, and the third, more briefly, is whether a graph neural network outperforms ordinary tabular models given the same features and blocks. The flood hazard is then combined with exposure and vulnerability on a 500 m grid, and the maps, data and code are released openly.

2 Related work

2.1 Flood susceptibility in Bangladesh

Table 1 summarises the studies closest to this one, each checked against its full text. They report high accuracies, from 0.81–0.87 for five models in the Teesta basin (Islam et al., 2021) to 0.984 for a random forest on the south-west coast (Adnan et al., 2023), but every study whose split could be read divided its points at random, and none holds out whole areas. The inventories also differ in kind, from points identified by field survey and local residents (Islam et al., 2021; Talukdar et al., 2020) or a water-board survey (Rudra and Sarkar, 2023) to pixels sampled from satellite flood maps (Islam and Chowdhury, 2024; Adnan et al., 2023), so the same accuracy figure can measure agreement with quite different things. Two studies compared the maps their models produced, and those comparisons are the more instructive. On the south-west coast four models each classified roughly 60 % of the land as flood-prone, yet their pixel-wise correlations ranged from 0.62 to 0.91 (Adnan et al., 2023). In one Sylhet sub-district a random forest and a support vector machine with AUCs of 0.964 and 0.954 agreed on only 73.3 % of the map (Islam and Chowdhury, 2024). Models that score equally well can therefore disagree about where flooding will happen, which is the property a map exists to convey. Multi-criteria hazard maps built from similar terrain factors, including TWI and HAND, reached validation scores of 0.77–0.79 in the Sangu basin and placed 39 % of roads and 25 % of education buildings in the highest hazard class (Zzaman et al., 2021), which shows why the exposure of infrastructure matters to planners.

2.2 Labels and validation

Radar offers training labels that cover every pixel, and Sentinel-1 flood maps have reached overall accuracies of 94.0–96.1 % (Twele et al., 2016) and 96.44 % for the 2017 Bangladesh floods (Uddin et al., 2019). Daily MODIS imagery has mapped 913 large floods worldwide (Tellman et al., 2021), although radar still struggles in vegetated and urban areas (Amitrano et al., 2024), and Sentinel-1 flood maps have already been used to estimate flood exposure across a whole country (Sherpa and Shirzaei, 2022). One Bangladesh study has already built its inventory from Sentinel-1, for a single Sylhet flood and with a random split (Islam and Chowdhury, 2024). Where no inventory exists, labels are often thresholded from the topographic wetness index (Beven and Kirkby, 1979) and height above nearest drainage (Nobre et al., 2011), the same indices that serve as predictors, so a model scored against such labels can succeed by relearning the rule that made them.

Random cross-validation of spatially structured data underestimates prediction error, and splitting by area can address this (Roberts et al., 2017). Splitting by time, training on past events and testing on later ones, has long been standard for landslides (Chung and Fabbri, 2003; Guzzetti et al., 2006), but among the Bangladesh flood studies only one map has been tested on later floods, a multi-criteria map of Sylhet scored against the floods of June 2022 and June 2024 (Swarup and Nury, 2026), and none has been set against the record of earlier floods. Where records and models have been compared, it was as hazard maps, with satellite records showing less frequent flooding than global flood models predict (Hawker

et al., 2020). Such records already exist for Bangladesh, as a national Sentinel-1 flood-frequency map (Haghighi, 2022), a weekly inundation series built for flood insurance (Thomas et al., 2023) and a two-decade reconstruction (Giezenanner et al., 2023). The effect can be large, since a forest-biomass model that explained more than half the variance under random validation showed almost no skill under spatial validation (Ploton et al., 2020). Blocking can force a model to extrapolate (Roberts et al., 2017) and has no theoretical basis for estimating map accuracy (Wadoux et al., 2021), so blocked scores are read here as performance on unseen areas of similar terrain, not as design-based map accuracy. Calibration is treated apart from ranking, by isotonic regression (Zadrozny and Elkan, 2002; Niculescu-Mizil and Caruana, 2005), and none of the Bangladesh studies read in full reports calibrated probabilities.

2.3 Graphs, interpolation and composite risk

Comparisons of many learning models are common in susceptibility mapping, and among 14 models for flood-prone areas of Pakistan boosted trees ranked first (Waleed and Sajjad, 2025). Graph neural networks learn from a node's neighbours and, in the inductive form used here, generalise to nodes unseen in training (Hamilton et al., 2017). In hazard mapping they appeared first for landslides (Wang et al., 2023a; Ma et al., 2025) and have since been applied to flood susceptibility on graphs of urban cells (Wang et al., 2023b) and of catchments linked by their topology (Liu et al., 2026), and a preprint reported a gain of 0.097 in AUC over the best tabular model on sub-watersheds linked by river connectivity (Sharma and Sharma, 2026). A review of 58 deep-learning flood studies found convolutional models usually the more accurate (Bentivoglio et al., 2022), and the searches made for this review found no graph model of the flood susceptibility of mapped infrastructure, nor one that links assets by proximity rather than by the flow of water. Composite indices such as INFORM combine hazard, exposure and vulnerability by geometric mean (European Commission Joint Research Centre, 2017), and the same aggregation is used here, with exposure counted asset by asset as Koks et al. (2019) did globally. That exposure depends on OpenStreetMap, whose building completeness exceeds 80 % in some cities but falls below 20 % in cities holding 48 % of the world's urban population (Herfort et al., 2023), and whose completeness across rural Bangladesh has not been measured.

3 Study area and data

Three regions with contrasting flood regimes are studied, riverine flooding in Rangpur and Rajshahi, flash flooding in the Sylhet haor basin and cyclone surge on the south-west coast (Fig. 1a, Table 2), and they share one set of model settings so that results stay comparable. Rangpur and Rajshahi, in the north-west, lie on the floodplains of the Teesta and the Brahmaputra–Jamuna, where monsoon rivers rise above their danger levels and stay there for days to weeks, and all five mapped floods fell between late June and October (Table 4). The Sylhet region holds the haor basin, a low bowl of seasonal wetlands fed by rivers from the hills to the north and east, where flash floods rise quickly and can arrive before the monsoon, as in April 2017, when

Study	Region	Model	Flood locations and split	Reported accuracy
Talukdar et al. (2020)	Teesta basin	bagging ensembles	207 flood, 206 non-flood points from inundation maps, topographic maps and local survey, random 80:20	AUC 0.945 (BgM5P)
Islam et al. (2021)	Teesta basin	Dagging, RS, RF, ANN, SVM	167 flood, 167 non-flood points from records, fieldwork and Google Earth, random 80:20	AUC 0.81–0.87 (validation)
Islam and Chowdhury (2024)	Sylhet, one sub-district	RF, SVM	1,500 flooded, 1,500 dry pixels from Sentinel-1, one event (June 2022), random 70:30	AUC 0.964 (RF)
Rudra and Sarkar (2023)	Sylhet division	ANN	1,280 flood and non-flood points from a Bangladesh Water Board survey, random 75:25	prediction rate 89–92 %
Swarup and Nury (2026)	Sylhet division	AHP, 19 indicators	independent flood inventories of June 2022 and June 2024, no training data	AUC 0.836 (2022), 0.794 (2024)
Adnan et al. (2023)	south-west coast	RF, KNN, MLP, GA-RBF-SVR	965 points sampled from a satellite flood-frequency map (1988–2012), 70:30 split	AUC 0.984 (RF), inter-model r 0.62–0.91
This system	Rangpur & Rajshahi	Gradient boosting	Sentinel-1, 5 events, 10 km blocks	mean AUC 0.817
This system	Sylhet	Gradient boosting	Sentinel-1, 4 events, 10 km blocks	mean AUC 0.857
This system	south-west coast	Gradient boosting	Sentinel-1, 4 cyclones, 10 km blocks	mean AUC 0.839

Table 1. Flood susceptibility studies in Bangladesh compared with this system. Every study was read in full except [Swarup and Nury \(2026\)](#), whose entries come from its abstract, and where a figure is given for training and validation separately the validation figure is shown. This system’s figures are means over 20 assignments of spatially held-out blocks, scored against observed floods, so the columns are not directly comparable, for the reasons given in the text.

they struck before the rice harvest. The south-west coast is a delta of tidal channels and polders, embanked islands of farmland, where flooding follows the landfall of a cyclone and the overtopping or breach of embankments rather than the stage of a river. The regions hold 23,016, 13,374 and 29,545 mapped assets, across parts of 21, 10 and 16 districts. Every dataset the pipeline ingests is openly licensed (Table 3). Administrative boundaries come from geoBoundaries, which unlike GADM may be redistributed and carries the union level for Bangladesh. OpenStreetMap coverage is uneven, so an area with few mapped assets shows low exposure whether or not much stands there, and WorldPop’s constrained product carries population only where buildings were detected in imagery.

4 Methods

The pipeline runs as a chain of steps (Fig. 2), and this section follows it in order. Settings not given here are in the Supporting Information (Section S1).

4.1 Assets, terrain and features

Assets are fetched from OpenStreetMap in three tag groups, lifelines (hospitals, schools, shelters, bridges, embankments), transport (primary to trunk roads, railways) and agriculture (farmland, aquaculture, canals, markets), and each is classified into one asset type. The configured boxes cross Bangladesh’s land border, so assets in India, or on land more than about 1 km outside the geoBoundaries border, are dropped, while those on chars, coastal ground and offshore are kept. Grid cells and grid-level validation are restricted the same way. From the 30 m SRTM elevation model, reprojected to UTM zone 46N, the pipeline derives slope, the topographic wetness index (TWI), height above nearest drainage (HAND) and D8 flow accumulation. FABDEM, which removes forest and building

heights from a 30 m elevation model ([Hawker et al., 2022](#)), would suit a flat delta better, but its licence forbids commercial use and requires derived data to carry the same terms, which the CC BY 4.0 data archive of this study could not. The vertical error of SRTM on flat floodplains is therefore one possible reason why TWI and HAND add little once elevation is known (Section 5.7), a reading that was not tested. Each asset carries these five terrain values at its representative point, WorldPop population density, its distance to the nearest hospital, school, flood shelter, road and irrigation channel, and its asset type, all standardised, and the pipeline stops if any feature is constant or the labels hold a single class.

4.2 Flood extents from Sentinel-1

Flood extents mapped from Sentinel-1 radar supply both the training labels and the reference against which the models are scored. Following the change-detection method recommended by UN-SPIDER ([UN-SPIDER, 2019](#)), an approach whose best Sentinel-1 variant reached an overall accuracy of 86.3 % against reference flood maps in an evaluation for Thailand ([Dadhich et al., 2026](#)), the minimum VV backscatter in each flood window is compared in Google Earth Engine with the median of a dry-season baseline from the same year. A pixel is flagged as flooded where it is below –16 dB and at least 3 dB darker than its baseline. Permanent water from the JRC record ([Pekel et al., 2016](#)), slopes above 5° and specks under eight pixels are removed, and the event masks are averaged into a flood-frequency layer (Fig. 1b–d). Five events are mapped in Rangpur and Rajshahi, four in Sylhet and four cyclone landfalls on the coast, each window set from ReliefWeb situation reports (Table 4).

The masks were checked against an independent map of two of the same floods, the August 2017 flood in Rangpur and Rajshahi and the April 2017 haor flood in Sylhet, which the




Figure 1. (a) The three study regions, drawn as their configured bounding boxes over the district boundaries of geoBoundaries, with neighbouring countries from Natural Earth. (b–d) Share of the mapped Sentinel-1 events in which each 100 m cell was flooded. Cells never seen flooded are left blank, and the removal of permanent water leaves the main river channels empty. Straight-edged bands follow the edges of radar swaths.

Region	Hazard	Bounding box (W, S, E, N)	Configuration
Rangpur & Rajshahi	riverine flood	88.0, 24.0, 89.9, 26.7	config.yaml
Sylhet	flash flood	90.6, 24.0, 92.7, 25.5	configs/sylhet.yaml
South-west coast	coastal and tidal flooding	88.9, 21.5, 90.8, 23.2	configs/sw_coastal.yaml

Table 2. Configured study regions. Each box is restricted to Bangladesh, and all use UTM zone 46N (EPSG:32646) for distance and area calculations.

Dataset	Use	Resolution	Licence
OpenStreetMap	infrastructure assets (roads, bridges, schools, hospitals, shelters, embankments, cropland)	vector	ODbL 1.0
SRTM 1 arc-second	elevation, slope, topographic wetness, height above drainage, flow accumulation	30 m	public domain
JRC Global Surface Water	historical water occurrence, one proxy-label vote	30 m	free with attribution
WorldPop constrained 2020	population density for vulnerability and population exposure	100 m	CC BY 4.0
geoBoundaries gbOpen	districts (ADM2), upazilas (ADM3), unions (ADM4)	vector	CC BY 4.0
Sentinel-1 GRD (via Earth Engine)	observed flood extents for training labels and validation	10 m, exported at 100 m	Copernicus open

Table 3. Datasets ingested by the pipeline. National files (WorldPop, geoBoundaries, JRC tiles) are downloaded once into a shared cache and clipped per region.




Figure 2. The processing chain, top to bottom. Source counts are totals over the three flood regions. Each asset enters the models as a vector of terrain, population, access-distance and asset-type features, with a label from the Sentinel-1 flood frequency. The asset model is expanded, with the graph network it replaced shown as the alternative it now is. Both models train on the same spatial blocks, and dashed connectors mark what the validation scores on the test blocks.

Global Flood Database maps from MODIS at 250 m (Tellman et al., 2021). On that grid, over cells that MODIS saw clear and that are not permanent water, the maps agree on 81.9 % and 82.0 % of cells, a figure dominated by dry ground. On flooded ground the agreement is only moderate, with a critical success index of 0.290 and 0.339 and a kappa of 0.346 and 0.406. The radar marks more ground as flooded than MODIS in Rangpur and Rajshahi (19.7 % against 13.2 %), consistent with flooded patches smaller than a MODIS pixel and with flooding hidden by cloud. It marks less in the Sylhet haor (13.2 % against 23.3 %), consistent with the difficulty radar has where water stands under vegetation (Amitrano et al., 2024), as it did under unharvested rice in April 2017. Neither explanation was tested, and neither map is ground truth.

4.3 Training labels

In the riverine regions ground counts as flood-prone where it flooded in at least two of the mapped events, while on the coast one cyclone window is enough, because surge drains within days and the six-to-twelve-day revisit leaves too few pixels flooded twice. The pipeline also writes a terrain-threshold ensemble, which marks a pixel as flood-prone where at least two of three indicators agree, TWI above 8, HAND below 5 m and JRC water occurrence above 25 %. TWI and HAND are also model inputs, so a model trained and scored on these labels can succeed by relearning the threshold rule.

4.4 Asset model

The default asset model is class-weighted gradient boosting with 300 iterations and a learning rate of 0.1, chosen by the comparison in Section 5.6, where it is set against a two-layer GraphSAGE network (Hamilton et al., 2017) on a six-nearest-neighbour graph of assets, a random forest and logistic regression. Validation holds out whole 10 km blocks, assigned at random to training (80 %), calibration and test, because a random split leaves every test asset beside training assets with almost the same features and label. Scores are calibrated by isotonic regression on the calibration blocks, with each level shrunk towards the middle when it holds few points, and every metric comes from the test blocks. In the two riverine regions the published map scores come from the same model given one more feature, the share of earlier events in which the ground flooded (Section 5.4). Every validation figure outside that section comes from the model without it, because the feature is built from the events the models are scored against.

4.5 Hazard surfaces

Of the two hazard surfaces that cover the grid, the first interpolates the asset scores by ordinary Kriging, which can add little between assets where the scores lack spatial structure. Its fitted variograms leave 0.307 of the sill unexplained at zero distance in Rangpur and Rajshahi and 0.742 in Sylhet. The second, the default, is a class-weighted logistic regression on each cell's own elevation, slope, TWI, HAND and flow accumulation, trained at the assets on the observed labels and blocks and evaluated at every 30 m cell. It leaves out the access distances, which describe exposure rather than the likelihood of flooding, and so follows individual channels and levees where the Kriged surface varies over tens of kilometres (Supporting Information, Fig. S1).

4.6 Exposure, vulnerability and composite risk

Exposure is measured twice for each 500 m cell, once as the type-weighted count of mapped assets and once as WorldPop density, both log-scaled. Vulnerability combines population density with the distance to the nearest hospital, flood shelter and primary road, and components without data are left out rather than filled with a constant. Composite risk is the geometric mean of hazard, exposure and vulnerability, as in INFORM (European Commission Joint Research Centre, 2017), so that no exposure means no risk, and cells are placed in five relative classes by quintile. Cells and administrative units with no mapped assets carry no score, since sparse mapping would otherwise read as low risk. Results are summarised for every union, upazila and district (Fig. 3), and the weights, the hotspot method and a sensitivity analysis are in the Supporting Information (Sections S2 and S3). The ranking of cells is robust to the chosen weights, with a mean Spearman correlation of 0.963 with the published ranking when every weight is perturbed at random in Rangpur and Rajshahi, and it changes materially only when vulnerability is reduced to population density alone ($\rho = 0.663$).

4.7 Validation design

Every model is scored on the held-out blocks by the area under the ROC curve (AUC), average precision and its lift over the base rate, and the Brier score, and the hazard surfaces are also scored against observed flooding on a grid of about 500 m cells, both over the whole grid and, refitted for each assignment, over the cells of its test blocks. Every comparison between models or label sources is repeated over 20 block assignments drawn with different seeds, because one assignment gives a single number with no sense of its spread. The paired differences are tested by the corrected resampled t-test of Nadeau and Bengio (2003), since the assignments reuse the same assets, and a plain t-test across them would overstate the evidence. A temporal test trains on the floods up to 2022 and scores against those of 2024 (Section 5.4). The 10 km blocks are narrower than the ranges of the variograms fitted to the asset scores, roughly 115, 14 and 31 km, so test assets near a block edge still share terrain with training assets, and the main results are repeated with 20 and 40 km blocks on the same seeds (Section 5.5).

5 Results

5.1 Held-out results

Table 5 reports one block assignment, the pipeline's default, which ranks 2, 1 and 17 of 20 by AUC in the three regions. In Sylhet 14.7 % of assets sit on flood-prone ground and the asset model reaches an AUC of 0.934. In Rangpur and Rajshahi only 6.7 % do, yet the model ranks well (AUC 0.884), with an average precision 7.0 times the base rate. Calibration lowers the Brier score from 0.069 to 0.029 in Rangpur and Rajshahi, slightly below the 0.032 of always predicting the base rate, and its deciles lie close to the diagonal in both riverine regions (Fig. 5). It carries less well to the coast, where the calibration blocks hold 5.5 % flood-prone assets against 1.7 % in the test blocks and the calibrated Brier score of 0.039 is worse than the base rate at 0.017.

Region	Window	Basis
Rangpur & Rajshahi	12–26 Aug 2017	second monsoon flood from 12 August; Dinajpur, Kurigram, Lalmonirhat and Rangpur among the worst affected
	12–26 Jul 2019	rain from 9 July, Brahmaputra–Jamuna rising to 18 July, Kurigram more than 1 m above danger level
	11–31 Jul 2020	second phase from 9–10 July, peak about 18 July, 27 basin gauges above danger into late July
	28 Jun–16 Jul 2024	northern floods from mid-June, Brahmaputra above danger at Nunkhawa, Hatia and Chilmari in early July
	27 Sep–12 Oct 2024	Teesta flooding of Lalmonirhat, Nilphamari, Kurigram and Rangpur from late September
Sylhet	30 Mar–20 Apr 2017	haor flooding from 28 March, striking Sunamganj on 1 April
	15–26 Jun 2022	heavy rain from 15 June, situation worst from 17 June
	16–28 Jun 2024	2.25 million affected in Sylhet, Sunamganj and Moulvibazar by 22 June
	18–30 Aug 2024	eastern floods from 20 August, including Moulvibazar and Habiganj
South-west coast	20–31 May 2020	Cyclone Amphan, landfall 20 May
	25 May–4 Jun 2021	Cyclone Yaas, landfall 26 May
	23 Oct–2 Nov 2022	Cyclone Sitrang, landfall 24–25 October
	26 May–5 Jun 2024	Cyclone Remal, landfall 26–27 May

Table 4. Flood events mapped from Sentinel-1. Windows bracket the flood peak; the baseline for each is the February–March dry season of the same year (mid-January to March for the pre-monsoon haor event). Sources are ReliefWeb situation reports, cited in the configuration files.




Figure 3. Mean infrastructure composite risk of each union in the three flood regions, in five relative classes by quintile of the unions scored in each region. Unions without mapped assets carry no score and are shaded as missing data, not as low risk. On the south-west coast the hazard term is close to chance against observed flooding (Section 5), so that panel reflects mapped exposure and vulnerability more than flood hazard.

Region	Assets	Positive rate	Asset AUC	Asset AP	Brier	Terrain AUC
Rangpur & Rajshahi	23,016	6.7 %	0.884	0.234	0.069	0.621
Sylhet	13,374	14.7 %	0.934	0.639	0.073	0.788
South-west coast	29,545	4.5 %	0.794	0.044	0.134	0.757

Table 5. Validation on held-out 10 km test blocks, for the pipeline’s default block assignment, from the model without the past-flooding feature that the riverine maps carry (Section 4.4). Positive rate is the share of assets on ground the radar saw flood in the required number of events. The Brier score is that of the raw class-weighted scores, and the calibrated probabilities are reported in the text. Means over repeated assignments, lower in the riverine regions and higher on the coast, are in Table 7.

5.2 Observed extents against terrain labels

The terrain-threshold ensemble captures 0.151 of the ground that flooded in at least 2 of the five Rangpur and Rajshahi events, with a false alarm ratio of 0.517, and marks 6.6 % of the region as flood-prone where the radar saw 21.2 % flood that often. Trained on those labels and scored on the same held-out blocks against observed flooding over 20 assignments, the default model reaches 0.716 against 0.817 for observed labels in Rangpur and Rajshahi, 0.770 against 0.857 in Sylhet and 0.758 against 0.839 on the coast. The gaps of 0.101, 0.087 and 0.080 favour the observed labels on 20, 19 and 19 of the assignments ($p < 0.001$, $p = 0.003$ and $p < 0.001$; Fig. 4). This test favours the observed labels, since that model is scored against the definition it learned, so Section 5.4 repeats it against later floods that neither label source was trained on.

5.3 The south-west coast

On the coast the asset model still ranks flooded ground above dry ground (AUC 0.794), but both hazard surfaces are weak. Over the whole grid, training ground included, the terrain surface reaches 0.542 and the Kriged one 0.538, against 0.657 and 0.666 in Rangpur and Rajshahi and 0.822 and 0.544 in Sylhet. Refitted for each assignment and scored only on the cells of its test blocks, which hold only ground near mapped assets, the terrain surface reaches 0.656, 0.763 and 0.610, so the whole-grid figure overstated it in Sylhet. Surge flooding is governed by polder embankments, tidal channels and landfall position, none of which elevation describes.

5.4 Transfer to later floods

The held-out blocks separate training from testing in space but not in time. Each model is therefore also trained on the floods up to 2022, keeping each region’s labelling rule capped by the number of earlier events, and scored on the held-out blocks against the ground any 2024 flood reached. Three references are scored on the same assets over 20 assignments (Table 6). The first is the share of earlier events in which the ground flooded, which is the map the radar record alone provides. The second is the same model trained with the 2024 events included, and the third is the same model trained on terrain-threshold labels, which draw on no flood event.

Trained only on earlier floods, the model ranks the 2024 flooding at 0.801 in Rangpur and Rajshahi, 0.797 in Sylhet and 0.807 on the coast. In both riverine regions the record of earlier flooding ranks it better, at 0.896 and 0.851, with the model ahead on only 2 and 4 of 20 assignments, whereas on the coast the model leads, at 0.807 against 0.668. The regions differ in how often floods return to the same ground, since only 16.9 % and 22.3 % of the assets flooded in 2024 stood on ground no earlier event had reached, against 65.5 % on the

coast, where each surge floods the stretch of coast where it lands. The 2024 floods in Rangpur and Rajshahi were also small, covering 4.9–8.3 % of the region against 19.9–24.8 % for each earlier flood, whereas in Sylhet they were as large as earlier ones.

Seven checks test whether the comparison favours one side (Supporting Information, Section S5, Table S1), and in Rangpur and Rajshahi the record keeps its lead under all of them, if narrowly under two. Sentinel-1B failed in December 2021, so the earlier events were seen by two satellites and the 2024 ones by one, and since each mask takes the lowest backscatter in its window, more scenes can map more flooding. Re-mapped from Sentinel-1A alone, the earlier events shrank by up to 15.7 % and the record still led by 0.097. The published model learns a yes-or-no label from the earlier events while the record keeps their fraction, and retrained on flooding in any earlier event, which is how the 2024 floods are counted, or on the fraction itself, the model fell further behind, by 0.133 and 0.131. With 40 km blocks the lead grew to 0.126, and setting aside the ground that radar could mistake for water every monsoon, such as wet paddy, narrowed it to within noise without reversing it. In Sylhet the record’s lead closed, to 0.017 and 0.016, when the model learned the same target as the test, and outside the haor’s seasonal water the two were level, so the Sylhet result rests on the choice of target. On the coast the model’s lead fell within noise with 20 and 40 km blocks.

The record can also enter the model as a feature, provided it comes from floods before the ones the model learns to predict, or the model would copy its own labels. The model is therefore trained to predict the last event year up to 2022 from its usual features and the share of the earlier events in which the ground flooded, and scored against the 2024 floods with that share taken over every event to 2022. In Rangpur and Rajshahi the combination beats both of its parts, at 0.924 against 0.896 for the record alone ($p = 0.020$) and 0.790 for the model without the feature. Its gain comes on the ground no earlier flood reached, 87.7 % of the test assets, where the record gives every asset a score of zero and the combined model still ranks the 2024 flooding at 0.729. In Sylhet the combination, at 0.849, cannot be told apart from either the model without the feature at 0.837 ($p = 0.196$) or the record alone at 0.851 ($p = 0.881$). It still ranks the 2024 flooding at 0.777 on the 79.8 % of test assets that no earlier flood reached. On the coast, where the design leaves a single cyclone to train on, the combination at 0.648 cannot be told apart from the record alone ($p = 0.580$).

The 2024 floods also give the label comparison a test that favours neither label source. The observed labels still lead, at 0.801 against 0.732 in Rangpur and Rajshahi, 0.797 against




Figure 4. ROC curves on the held-out test blocks of each flood region, scored against the observed definition of flood-prone ground. Both curves come from the default gradient-boosted model on one block assignment, sharing features and blocks and differing only in their training labels. The dotted diagonal is a random ranking, and n is the number of assets in the test blocks.




Figure 5. Reliability of the flood model on the test blocks, with assets grouped into deciles of each prediction. The raw class-weighted score overstates the likelihood of flooding. The isotonic calibration, fitted on other blocks, brings the riverine deciles close to the diagonal. On the coast, where the calibration blocks hold more flood-prone ground than the test blocks (5.5 % against 1.7 %), it still overstates the flooding observed.

Region	Model to 2022	All events	Past flooding	Terrain labels
Rangpur & Rajshahi	0.801 $\pm$ 0.063	0.794	0.896 $\pm$ 0.042	0.732 $\pm$ 0.067
Sylhet	0.797 $\pm$ 0.064	0.850	0.851 $\pm$ 0.039	0.773 $\pm$ 0.043
South-west coast	0.807 $\pm$ 0.055	0.834	0.668 $\pm$ 0.072	0.736 $\pm$ 0.066

Table 6. AUC against the 2024 flood extents on held-out 10 km test blocks, as mean $\pm$ standard deviation over 20 block assignments. The model to 2022 is trained only on the earlier events, the all-events model on labels that include 2024, past flooding is the share of the earlier events in which the ground flooded, and the terrain-label model is the same model trained on the terrain-threshold labels. The 2024 events reached 2.8 %, 17.7 % and 2.9 % of assets in the three regions.

0.773 in Sylhet and 0.807 against 0.736 on the coast, ahead on 18, 15 and 18 of 20 assignments. The coastal gap of 0.071 stays close to its spatial 0.080, while the riverine gaps shrink, from 0.101 to 0.069 in Rangpur and Rajshahi and from 0.087 to 0.024 in Sylhet, where the remaining gap cannot be told apart from noise. Much of the Sylhet advantage in the spatial test therefore came from the shared definition rather than from the labels themselves.

5.5 Larger blocks

Wider blocks lower the held-out AUC of the default model only a little in the riverine regions, from 0.817 at 10 km to 0.791 at 40 km in Rangpur and Rajshahi and from 0.857 to 0.836 in Sylhet, and more on the coast, from 0.839 to 0.774, so the 10 km figures carry a small optimism, largest on the coast. The comparison with the flood record under wider blocks is reported in Section 5.4.

5.6 Model comparison

A graph neural network, the design this system began with, is set against logistic regression, a random forest and gradient boosting on the same features, assets and blocks, with the wetness index alone as a model-free reference, over 20 assignments (Table 7, Fig. 6). Gradient boosting scores best in every region, and the graph network trails it measurably only in Rangpur and Rajshahi, by 0.028 in AUC ($p = 0.003$), while the wetness index alone ranks close to chance, at 0.517, 0.535 and 0.577. Gradient boosting is therefore the default for every region, since choosing a model per region would fit the model to the split. Its AUC varies by 0.042 to 0.046 from one assignment to the next, so each single-assignment figure, including those in Table 5, is one draw from a distribution of that width.

5.7 What the model relies on

Shuffling each feature across the test assets over the same assignments shows what the default model uses (Fig. 7). Elevation carries almost all of the terrain signal, costing 0.086, 0.220 and 0.129 in AUC when shuffled, while TWI, HAND, slope and flow accumulation cost at most 0.000, 0.009 and 0.007. Shuffled together, since correlated features can stand in for one another, the five terrain features cost 0.080, 0.214 and 0.132, about as much as elevation alone, and slightly less in Rangpur and Rajshahi, a difference within the noise of the shuffling. The five access distances together cost 0.110, 0.133 and 0.142, more than the terrain in Rangpur and Rajshahi and on the coast and second to it in Sylhet. They describe distance to services rather than any property of flooding and most likely stand in for where roads and settlements sit on the floodplain, a reading that was not tested directly.

6 Discussion

6.1 What the flood record already knows

Susceptibility studies rarely ask whether a model anticipates a later flood better than the record of earlier ones, and here the answer followed how often floods returned to the same ground. In Rangpur and Rajshahi, where most of the ground the 2024 floods reached had flooded before, the record led under every check applied. The model added to it only when given the record as a feature, gaining on ground no earlier flood had reached, where the record cannot rank at all. In the

Sylhet haor the record's lead depended on what the model was trained to predict and lay on the margins of the seasonal water. A published multi-criteria map scored on the same 2024 floods reached 0.794 (Swarup and Nury, 2026), close to this model's 0.797, while the record ranked them at 0.851, a comparison that is only indicative, since that map was scored over areas rather than assets. On the coast, where most of the ground Cyclone Remal flooded had not flooded in the earlier cyclones, the model's lead did not survive wider blocks. Three regions cannot establish that pattern, but they show that the comparison has to be made region by region rather than assumed. The record marks where water returns and the terrain ranks where it has not yet been seen, so a susceptibility map for a river basin with a radar flood history should at least report the record beside it, and where the record is informative the two can be combined. A study that reports accuracy without that reference cannot show that its model knows more than the flood record already does, and the comparison is simple to make wherever Sentinel-1 has observed earlier floods.

6.2 Labels decide what an accuracy figure means

The label source decides what a published accuracy figure measures, since terrain-threshold labels cost the model 0.101, 0.087 and 0.080 in AUC on held-out blocks. Against later floods that neither label source had seen they still cost 0.069, 0.024 and 0.071, so in Rangpur and Rajshahi and on the coast the advantage of radar labels does not depend on the test sharing their definition, while in Sylhet most of it did. The indices those labels are built on, TWI and HAND, add almost nothing to the model once elevation is known, and a model scored against labels derived from its own inputs measures its recovery of the labelling rule rather than its skill at locating floods. The mean AUCs here, 0.817 and 0.857 in the riverine regions, sit at or below the best-model AUCs of 0.87 to 0.98 in Table 1. The comparison is only indicative, since those studies split their points at random while every figure here comes from 10 km blocks the model never saw, and wider blocks lower these figures a little further (Section 5.5). Blocking has its own critics (Roberts et al., 2017; Wadoux et al., 2021), and the blocked scores estimate performance on unseen areas of similar terrain rather than design-based map accuracy.

6.3 What the graph and interpolation did not add

On identical features and blocks the graph network trailed the best tabular model in every region, measurably only in Rangpur and Rajshahi, probably because the features already carry the neighbourhood, since nearby assets share elevation and distances to the same services, although this was not tested. The graphs that reported gains elsewhere linked catchments or sub-watersheds along the flow of water (Liu et al., 2026; Sharma and Sharma, 2026), whereas this one linked each asset to its nearest neighbours, so the result bears on proximity graphs of assets, and whether a river-connected graph would help remains untested. Kriging fared no better, and over the whole grid in Sylhet it ranked observed flooding at 0.544 against 0.822 for the terrain surface. Studies proposing either should report the tabular baseline on the same blocks.

6.4 Implications for flood risk management

First, where Sentinel-1 has observed several floods, the record of where water has already stood was at least as good a guide

Model	Rangpur & Rajshahi	Sylhet	South-west coast
Graph model (GraphSAGE)	0.789 ± 0.047	0.854 ± 0.039	0.817 ± 0.057
Gradient boosting	0.817 ± 0.045	0.857 ± 0.042	0.839 ± 0.046
Random forest	0.816 ± 0.044	0.845 ± 0.052	0.838 ± 0.049
Logistic regression	0.786 ± 0.039	0.855 ± 0.036	0.809 ± 0.046
Wetness index alone	0.517 ± 0.030	0.535 ± 0.029	0.577 ± 0.041
Graph minus best baseline	-0.028	-0.003	-0.022

Table 7. AUC on held-out test blocks, as mean $\pm$ standard deviation over 20 block assignments. Every model sees the same features at the same assets and the same blocks, so the only difference is the model. The last row is the paired difference between the graph network and whichever other model scores best in that region. Gradient boosting is the system’s default.




Figure 6. AUC of each model on the held-out test blocks, one circle per block assignment and a vertical bar at the mean. The graph network is in blue, the other models in orange, and the dotted line marks a random ranking. The graph network sits inside the range of the ordinary models in every region.




Figure 7. Loss in AUC on the held-out test blocks when each feature is shuffled across the test assets, as mean and standard deviation over 20 block assignments, for the default gradient-boosted model. The AUC in each title is the mean before shuffling.

to the next riverine flood as a terrain model alone, and in Rangpur and Rajshahi a better one. It costs little to compute, so agencies such as the Bangladesh Water Development Board and the Flood Forecasting and Warning Centre could publish it as a planning layer beside any susceptibility map. The terrain model adds most where the record is silent, on ground no mapped flood has reached and possibly on a coast where each surge strikes new ground, although the coastal lead did not survive wider blocks. Second, a susceptibility map used to site shelters, embankments or schools should be judged on held-out areas and later floods rather than on random splits, since the accuracy a random split reports overstates how far the map can be trusted in places it has not seen. Third, the composite risk classes rank unions within a region and must be read with their gaps, because a union without mapped assets carries no score rather than a low one. The best tabular models tested here were almost indistinguishable in AUC, so better OpenStreetMap coverage of rural unions would probably improve the map more than a new model would. Infrastructure also fails in networks, and the disruption of services beyond a flooded asset, which asset-level exposure omits, needs network modelling of the kind developed for Accra (Schotten and Bachmann, 2023). Validated against losses, maps of this kind could help set priorities among structural and non-structural flood measures in the districts they cover.

6.5 Limits of the maps

The coast is the clearest limit, since both hazard surfaces are weak there, the terrain surface reaching 0.610 on held-out ground, because surge flooding follows embankments, tidal channels and landfall position, which the pipeline does not yet ingest. The coastal composite is therefore not a map of surge exposure. Exposure depends on OpenStreetMap, whose rural completeness is unmeasured, and composite risk is defined only where assets are mapped, 12.5 % of grid cells in Rangpur and Rajshahi, while vulnerability omits income, housing and age structure, dimensions that a household survey in flood-prone Sunamganj found to matter (Rahman et al., 2025). The radar labels carry error of their own, agreeing only moderately with an independent MODIS map and counting inundated paddy as flooding. The temporal test rests on a single flood year, and in Rangpur and Rajshahi on small floods that stayed largely within earlier extents, so a larger flood reaching new ground might favour the model, as on the coast. Every score here therefore measures agreement with flooding mapped by radar rather than with flooding on the ground. The calibrated probabilities hold for a region as a whole rather than for any one area. The composite risk has not been checked against losses, and damage records for the 2022 and 2024 floods, where available, or depth-damage curves for infrastructure components such as those derived by expert elicitation (Williams et al., 2025), would allow the risk classes to be tested against the harm they are meant to anticipate. None of the surfaces is a forecast, and warnings remain the task of the Flood Forecasting and Warning Centre and the Bangladesh Meteorological Department.

7 Conclusions

In the riverine region where floods returned to the same ground, a terrain-based flood susceptibility model did not know more than the flood record. In Rangpur and Rajshahi the record of earlier floods predicted the 2024 flooding better than the model under every check applied, and a model given that record as a feature beat both, reaching 0.924, by ranking the ground no earlier flood had reached. In the Sylhet haor the comparison turned on the model's training target, and on the cyclone-struck coast it remains open. Radar-observed flood extents made better training labels than terrain thresholds, by 0.101, 0.087 and 0.080 in AUC on held-out ground and by 0.069, 0.024 and 0.071 against later floods. Ordinary tabular models matched or beat a graph neural network, and the terrain hazard surface ranked observed flooding on held-out ground at 0.656 and 0.763 in the riverine regions and only 0.610 on the coast. Every number in this paper is generated from the released pipeline outputs, so each can be checked and rebuilt.

Data availability statement

The data that support the findings of this study are openly available in Zenodo at <https://doi.org/10.5281/zenodo.23130732>. The code that produced them is openly available at <https://github.com/rakibhridoy/rime-s-mapathon> under the MIT licence. The repository holds the pipeline, region configurations, dashboard and test suite. The pipeline also holds a landslide susceptibility module for the Chittagong Hill Tracts, which is described in the technical documentation and will be reported separately. The Zenodo archive packages the input data and outputs for all four regions, including that one, per region with a manifest of checksums and licences. Every figure in this paper is written into the text by the analysis scripts in the code repository directly from the pipeline outputs, so a rerun of the pipeline updates text and results together. The results are published through an open dashboard at <https://fermium.systems/hazmapper>, which shows the risk maps, rankings and per-unit summaries without a sign-in and states the limitations of Section 6 beside the maps. An earlier technical report on this system, prepared for the RIMES Mapathon in March 2026, contained errors that were found in a later review, the most serious being a coordinate-system error that left every training label at zero. Each error and its correction is listed in the technical documentation in the code repository, and no result of that report is used here.

Declarations

Competing interests. Md Rakib Hasan is affiliated with Fermium Systems, which hosts the public dashboard described here. The software is released under the MIT licence and the data under the open licences listed in Table 3, and no result depends on a proprietary component.

Author contributions. Md Rakib Hasan is the sole author and carried out the conceptualisation, methodology, software, formal analysis, investigation, validation, data curation, visualisation and writing.

Ethics approval and consent. Not applicable. The study uses public datasets only and involves no human participants.

Acknowledgements

The author received no specific funding for this work. Infrastructure data are © OpenStreetMap contributors, under the Open Database Licence. Population data are from WorldPop, University of Southampton, CC BY 4.0. Administrative boundaries are from geoBoundaries, William & Mary geoLab, CC BY 4.0. Neighbouring country outlines in Figure 1 are from Natural Earth, in the public domain. Surface-water occurrence is from the Joint Research Centre of the European Commission (Pekel et al., 2016). Elevation is from the NASA Shuttle Radar Topography Mission. Sentinel-1 data are from the Copernicus programme of the European Union. The system was first developed for the RIMES Mapathon, ResilienceAI track.

References

- Adnan, M. S. G., Siam, Z. S., Kabir, I., Kabir, Z., Ahmed, M. R., Hassan, Q. K., Rahman, R. M., and Dewan, A. (2023). A novel framework for addressing uncertainties in machine learning-based geospatial approaches for flood prediction. *Journal of Environmental Management*, 326:116813.
- Amitrano, D., Di Martino, G., Di Simone, A., and Imperatore, P. (2024). Flood Detection with SAR: A Review of Techniques and Datasets. *Remote Sensing*, 16(4):656.
- Bentivoglio, R., Isufi, E., Jonkman, S. N., and Taormina, R. (2022). Deep learning methods for flood mapping: a review of existing applications and future research directions. *Hydrology and Earth System Sciences*, 26(16):4345–4378.
- Beven, K. J. and Kirkby, M. J. (1979). A physically based, variable contributing area model of basin hydrology / Un modèle à base physique de zone d'appel variable de l'hydrologie du bassin versant. *Hydrological Sciences Bulletin*, 24(1):43–69.
- Chung, C.-J. F. and Fabbri, A. G. (2003). Validation of spatial prediction models for landslide hazard mapping. *Natural Hazards*, 30:451–472.
- Dadhich, G., Babel, M., Miyazaki, H., Pal, I., Tripathi, N. K., and Virdis, S. G. P. (2026). Performance Evaluation of Automated Flood Inundation Mapping Techniques Using Multi-Temporal Sentinel-1 SAR Data. *Journal of Flood Risk Management*, 19(2):e70201.
- European Commission Joint Research Centre (2017). Index for Risk Management – INFORM: Concept and Methodology, Version 2017. Technical report, Publications Office of the European Union.
- Giezendanner, J., Mukherjee, R., Purri, M., Thomas, M., Mauerman, M., Islam, A. K. M. S., and Tellman, B. (2023). Inferring the past: a combined CNN–LSTM deep learning framework to fuse satellites for historical inundation mapping. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops*, pages 2155–2165.
- Guzzetti, F., Reichenbach, P., Ardizzone, F., Cardinali, M., and Galli, M. (2006). Estimating the quality of landslide susceptibility models. *Geomorphology*, 81:166–184.
- Haghighi, M. H. (2022). Large-scale mapping of flood using Sentinel-1 radar remote sensing. In *The International Archives of the Photogrammetry, Remote Sensing and Spatial Information Sciences*, volume XLIII-B3-2022, pages 1097–1102.
- Hamilton, W. L., Ying, R., and Leskovec, J. (2017). Inductive Representation Learning on Large Graphs. *Advances in Neural Information Processing Systems* 30; arXiv:1706.02216.
- Hawker, L., Neal, J., Tellman, B., Liang, J., Schumann, G., Doyle, C., Sullivan, J. A., Savage, J., and Tshimanga, R. (2020). Comparing earth observation and inundation models to map flood hazards. *Environmental Research Letters*, 15:124032.
- Hawker, L., Uhe, P., Paulo, L., Sosa, J., Savage, J., Sampson, C., and Neal, J. (2022). A 30 m global map of elevation with forests and buildings removed. *Environmental Research Letters*, 17(2):024016.
- Herfort, B., Lautenbach, S., Porto de Albuquerque, J., Anderson, J., and Zipf, A. (2023). A spatio-temporal analysis investigating completeness and inequalities of global urban building data in OpenStreetMap. *Nature Communications*, 14(1):3985.
- Islam, A. R. M. T., Talukdar, S., Mahato, S., Kundu, S., Eibek, K. U., Pham, Q. B., Kuriqi, A., and Linh, N. T. T. (2021). Flood susceptibility modelling using advanced ensemble machine learning models. *Geoscience Frontiers*, 12(3):101075.
- Islam, R. and Chowdhury, P. (2024). Local-scale flash flood susceptibility assessment in northeastern Bangladesh using machine learning algorithms. *Environmental Challenges*, 14:100833.
- Koks, E. E., Rozenberg, J., Zorn, C., Tariverdi, M., Voudoukas, M., Fraser, S. A., Hall, J. W., and Hallegatte, S. (2019). A global multi-hazard risk analysis of road and railway infrastructure assets. *Nature Communications*, 10(1):2677.
- Liu, J., Zhao, G., Xiong, J., and Kinouchi, T. (2026). A graph-based deep learning approach for daily flash flood susceptibility modeling in China. *Water Resources Research*, 62:e2025WR041360.
- Ma, J., Han, Z., Liu, F., Wang, X., Hu, J., and Zhang, P. (2025). ConToGCN: A landslide susceptibility assessment model considering contour topographic features in slope units using graph convolution network. *CATENA*, 255:109029.
- Nadeau, C. and Bengio, Y. (2003). Inference for the generalization error. *Machine Learning*, 52(3):239–281.
- Niculescu-Mizil, A. and Caruana, R. (2005). Predicting good probabilities with supervised learning. In *Proceedings of the 22nd international conference on Machine learning - ICML '05*, pages 625–632.
- Nobre, A., Cuartas, L., Hodnett, M., Rennó, C., Rodrigues, G., Silveira, A., Waterloo, M., and Saleska, S. (2011). Height Above the Nearest Drainage – a hydrologically relevant new terrain model. *Journal of Hydrology*, 404(1-2):13–29.

- Pekel, J.-F., Cottam, A., Gorelick, N., and Belward, A. S. (2016). High-resolution mapping of global surface water and its long-term changes. *Nature*, 540(7633):418–422.
- Ploton, P., Mortier, F., Réjou-Méchain, M., Barbier, N., Picard, N., Rossi, V., Dormann, C., Cornu, G., Viennois, G., Bayol, N., Lyapustin, A., Gourlet-Fleury, S., and Pélissier, R. (2020). Spatial validation reveals poor predictive performance of large-scale ecological mapping models. *Nature Communications*, 11(1):4540.
- Rahman, M. M., Tanni, K. N., Shobuj, I. A., Hossain, M. T., Alam, E., Al Hattaw, K. S., and Islam, M. K. (2025). Multidimensional Vulnerability Assessment for Floods: Evidence From Flood-Prone Areas of Bangladesh. *Journal of Flood Risk Management*, 18(3):e70089.
- Roberts, D. R., Bahn, V., Ciuti, S., Boyce, M. S., Elith, J., Guillerá-Arroita, G., Hauenstein, S., Lahoz-Monfort, J. J., Schröder, B., Thuiller, W., Warton, D. I., Wintle, B. A., Hartig, F., and Dormann, C. F. (2017). Cross-validation strategies for data with temporal, spatial, hierarchical, or phylogenetic structure. *Ecography*, 40(8):913–929.
- Rudra, R. R. and Sarkar, S. K. (2023). Artificial neural network for flood susceptibility mapping in Bangladesh. *Heliyon*, 9(6):e16459.
- Schotten, R. and Bachmann, D. (2023). Critical infrastructure network modelling for flood risk analyses: Approach and proof of concept in Accra, Ghana. *Journal of Flood Risk Management*, 16(3):e12913.
- Sharma, P. and Sharma, S. (2026). Flood risk follows valleys, not grids: Graph neural networks for flash flood susceptibility mapping in Himachal Pradesh with conformal uncertainty quantification. arXiv preprint 2603.15681, not peer reviewed.
- Sherpa, S. F. and Shirzaei, M. (2022). Country-wide flood exposure analysis using Sentinel-1 synthetic aperture radar data: Case study of 2019 Iran flood. *Journal of Flood Risk Management*, 15(1):e12770.
- Swarup, S. D. and Nury, A. H. (2026). Flood risk prioritization in data-scarce regions: a tri-dimensional GIS-AHP framework for northeastern Bangladesh. *Journal of Water and Climate Change*, 17:1229–1256.
- Talukdar, S., Ghose, B., Shahfahad, Salam, R., Mahato, S., Pham, Q. B., Linh, N. T. T., Costache, R., and Avand, M. (2020). Flood susceptibility modeling in Teesta River basin, Bangladesh using novel ensembles of bagging algorithms. *Stochastic Environmental Research and Risk Assessment*, 34(12):2277–2300.
- Tellman, B., Sullivan, J. A., Kuhn, C., Kettner, A. J., Doyle, C. S., Brakenridge, G. R., Erickson, T. A., and Slayback, D. A. (2021). Satellite imaging reveals increased proportion of population exposed to floods. *Nature*, 596(7870):80–86.
- Thomas, M., Tellman, E., Osgood, D., DeVries, B., Islam, A., Steckler, M., Goodman, M., and Billah, M. (2023). A framework to assess remote sensing algorithms for satellite-based flood index insurance. *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, 16:2589–2604.
- Twele, A., Cao, W., Plank, S., and Martinis, S. (2016). Sentinel-1-based flood mapping: a fully automated processing chain. *International Journal of Remote Sensing*, 37(13):2990–3004.
- Uddin, K., Matin, M. A., and Meyer, F. J. (2019). Operational Flood Mapping Using Multi-Temporal Sentinel-1 SAR Images: A Case Study from Bangladesh. *Remote Sensing*, 11(13):1581.
- UN-SPIDER (2019). Recommended Practice: Flood Mapping and Damage Assessment Using Sentinel-1 SAR Data in Google Earth Engine. <https://un-spider.org/advisory-support/recommended-practices/recommended-practice-google-earth-engine-flood-mapping>.
- Wadoux, A. M.-C., Heuvelink, G. B., de Bruin, S., and Brus, D. J. (2021). Spatial cross-validation is not the right way to evaluate map accuracy. *Ecological Modelling*, 457:109692.
- Waleed, M. and Sajjad, M. (2025). Advancing flood susceptibility prediction: A comparative assessment and scalability analysis of machine learning algorithms via artificial intelligence in high-risk regions of Pakistan. *Journal of Flood Risk Management*, 18(1):e13047.
- Wang, X., Du, A., Hu, F., Liu, Z., Zhang, X., Wang, L., and Guo, H. (2023a). Landslide susceptibility evaluation based on active deformation and graph convolutional network algorithm. *Frontiers in Earth Science*, 11:1132722.
- Wang, Z., Lyu, H., and Zhang, C. (2023b). Pluvial flood susceptibility mapping for data-scarce urban areas using graph attention network and basic flood conditioning factors. *Geocarto International*, 38:2275692.
- Williams, J. H., Vu, H. N., Paulik, R., Zorn, C., and Wotherpoon, L. (2025). Expert-Derived Flood Damage Curves for Critical Infrastructure Network Components. *Journal of Flood Risk Management*, 18(2):e70045.
- Zadrozny, B. and Elkan, C. (2002). Transforming classifier scores into accurate multiclass probability estimates. In *Proceedings of the eighth ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 694–699.
- Zzaman, R. U., Nowreen, S., Billah, M., and Islam, A. S. (2021). Flood hazard mapping of Sangu River basin in Bangladesh using multi-criteria analysis of hydrogeomorphological factors. *Journal of Flood Risk Management*, 14(3):e12715.